Full-Spectrum Leadership Coach — Product Requirements Document

Version: 3.0 (MVP-ready, scalability-first) **Date:** 20 April 2026 **Status:** Draft for Engineering
Source documents: `Full_Spectrum_Leadership_Coach_PRD_v2.pdf`, `innergy_whatsapp_ux_flow_docx.pdf`, `FLS_Diagnostic_for_CXOs_docx.pdf`

---

0. How to read this document

Section 1–4 are the product frame. Section 5 is the complete functional spec. Section 6 is the data-driven configuration model (this is where "nothing is hardcoded" is operationalised). Section 7 is the full V1 diagnostic content (every question, option, score, and band) as it will be seeded into the database on day one. Sections 8–11 cover non-functional requirements, acceptance criteria, and open items.

If you are building the MVP, **Sections 5, 6, 7, and 10** are the ones you will live in.

---

1. Product summary

The Full-Spectrum Leadership Coach is a **multi-tenant, WhatsApp-delivered AI platform** that helps senior leaders understand and improve their leadership readiness across three dimensions: **Cognitive Clarity**, **Relational Influence**, and **Inner Mastery**.

The long-term product (per PRD v2) is a full AI coach: it diagnoses, coaches in real time, grounds its responses in company-specific context (values, strategy, frameworks), supports voice, and escalates safely.

V1 (this MVP) delivers the diagnostic and handoff layer of that product — branded for Innergy as the *AI Leadership Readiness Assessment* — and builds the data model, tenancy model, and content pipeline the full coach will later sit on top of. V1 ends each user journey with a human CTA to the tenant's assigned coach (for Innergy, Rashmi Sharma).

2. Goals and non-goals

2.1 Goals (V1)

- Deliver a fully working diagnostic + results experience on WhatsApp for multiple tenants from day one.
- Make every piece of content — questions, options, scores, bands, interpretation copy, tenant branding, coach details — editable through an admin UI without a code deploy.
- Support both **templated interpretation** and **LLM-generated per-user interpretation**, switchable per tenant.
- Support voice input and voice output throughout the flow as a user-selectable mode.
- Architect the system so that real-time AI coaching, memory, and knowledge-base grounding (per PRD v2) can be added in later phases without re-platforming.

2.2 Non-goals (V1)

- Real-time AI coaching conversations (V2).
- Manager / 360 perception-gap assessments (V2+, per PRD v2 §19).
- Native mobile or web chat (V1 is WhatsApp only).
- Self-serve tenant onboarding (V1 onboards tenants manually via internal admin).

3. Phasing

Phase	Scope	Status
V1 (MVP)	Multi-tenant diagnostic bot on WhatsApp · voice I/O · scoring · interpretation (templated + LLM) · results visual · human coach CTA · admin dashboard · analytics	This PRD
V2	Real-time AI coaching on lowest dimension · knowledge-base grounding (tenant values, strategy, frameworks) · longitudinal user memory · escalation	Deferred

Phase	Scope	Status
	engine	
V3+	Manager/360 scoring with perception-gap insights · team aggregation · cross-tenant benchmarking (opt-in)	Deferred

4. Users

Role	Description	V1 capabilities
Leader (end user)	CXO / senior manager at a tenant company	Take diagnostic, receive results, book coach
Coach	Named human coach assigned to a tenant (e.g. Rashmi for Innergy)	Receives warm-lead notifications, calendar bookings
Tenant admin	HR, L&D, or consultant inside the tenant	View their tenant's analytics, manage users, upload company context (knowledge base, for V2)
Super admin (Innergy ops)	Internal platform operator	Onboard new tenants, edit global question bank, edit per-tenant overrides, manage coaches, view cross-tenant analytics

5. Functional requirements

Numbered FRs are the atomic requirements for implementation. Each is tagged `[V1]` or `[V2+]`. Only `[V1]` items are in scope for this build.

5.1 Tenancy and onboarding

- **FR-1.1 [V1]** Every domain record (user, session, answer, result, content item, notification) must carry a `tenant_id`. No cross-tenant reads are permitted except by a super admin.
- **FR-1.2 [V1]** Each tenant has a configuration record containing: display name, logo URL, brand colours (primary/secondary), coach name + coach booking URL + coach notification channel, assigned WhatsApp sender number (or a shared number with tenant routing — see FR-1.4), LinkedIn URL, closing-message text, and feature flags (voice enabled, LLM interpretation enabled, email capture enabled).
- **FR-1.3 [V1]** Each tenant has an **instrument** assignment — a specific version of the diagnostic (questions + options + scores + bands). Multiple tenants can share one instrument, or a tenant can have its own.
- **FR-1.4 [V1] Tenant routing on WhatsApp** — two supported modes:
  - *Dedicated number mode*: tenant has its own WhatsApp Business number; tenant is inferred from the inbound number.
  - *Shared number mode*: a single WhatsApp number is used; tenant is inferred from the QR-code deep-link trigger payload (e.g. `START_innergy`, `START_acme`). Super admin generates one QR per tenant.
- **FR-1.5 [V1]** Super admin can create, edit, and soft-disable tenants from the admin dashboard. Soft-disable stops new sessions but preserves history.

5.2 Entry and welcome

- **FR-2.1 [V1]** Entry point is a QR code scan that opens WhatsApp pre-populated with a tenant-specific trigger payload. The bot must activate on receipt without requiring the user to type first.
- **FR-2.2 [V1]** The welcome sequence is **3 messages**, each pulled from tenant-configurable templates with variable substitution (`{{coach_name}}`, `{{tenant_name}}`, `{{duration}}`, `{{dimension_names}}`). Default copy matches `innergy_whatsapp_ux_flow` §STEP 2.

- **FR-2.3 [V1]** The welcome ends with a `YES / LATER` choice. On `LATER`, schedule one reminder 24h later, and at most one more 24h after that. No further nudges.
- **FR-2.4 [V1]** After `YES`, capture **first name** and **organisation** as two sequential prompts. Email is **not** requested at this stage.
- **FR-2.5 [V1]** Name and organisation are stored on the user record and used for personalisation. Organisation is a free-text field in V1 (not validated against the tenant record).

5.3 Diagnostic delivery

- **FR-3.1 [V1]** The diagnostic is delivered **one question at a time**. The bot must not send the next question until a valid answer is received for the current one.
- **FR-3.2 [V1]** Each question message includes a progress indicator (`Q3 of 8`) and the section name. The question text, options, and progress-indicator format are all template-driven (§6.3).
- **FR-3.3 [V1]** The user answers by typing `A`, `B`, `C`, or `D` (case-insensitive, whitespace-trimmed). Additional accepted forms: `a)`, `A.`, `option a`. These normalisation rules are configurable.
- **FR-3.4 [V1]** On invalid input, respond with the tenant-configured invalid-answer template. Do not advance. Log the invalid attempt.
- **FR-3.5 [V1]** Section transitions send one configurable message acknowledging completion and introducing the next section. The transition into Relational Influence must explicitly frame the lens as *"about the team, not the individual"* (the instrument text controls this; it is not hardcoded).
- **FR-3.6 [V1]** After the final question, display a completion acknowledgement and a deliberate **2–3 second delay** before results, to create a sense of calculation.
- **FR-3.7 [V1] Abandonment handling**: if the user does not reply for 24h mid-session, send one reminder referencing the exact question they were on. If no reply after a further 48h, close the session gracefully; require a fresh QR scan to restart. Reminder copy and timings are tenant-configurable.

5.4 Scoring engine

- **FR-4.1 [V1]** Each option on each question carries a score value (integer). On answer, record the chosen option, its score, and a timestamp.
- **FR-4.2 [V1]** Section score = sum of option scores for that section. Overall score = sum of section scores.
- **FR-4.3 [V1]** Band for a section score is derived from a configurable lookup: `{section_id, min_score, max_score, band_label}`. Overall band derives from its own configurable lookup.
- **FR-4.4 [V1]** Scoring must be **pure** — given the answers and the instrument version, the result is deterministic and reproducible. Retroactive band edits must not mutate historical result records; they create a new interpretation read on replay.
- **FR-4.5 [V1]** The system must store the instrument version ID on every result so historical results remain readable after the instrument is edited.

5.5 Interpretation engine

- **FR-5.1 [V1]** Interpretation runs in one of two modes, set per tenant (and overridable per instrument version):
  - **Template mode**: for each `(dimension, band)` pair, a tenant-configured interpretation string is selected. Variables (`{{name}}`, `{{organisation}}`, `{{score}}`, `{{max_score}}`, `{{lowest_dimension}}`) are interpolated.
  - **LLM mode**: a prompt template is filled with the user's scores, bands, name, organisation, the tenant's company-context document (V2+), and any instrument-level interpretation guidance, then sent to the configured LLM. The response is returned as the interpretation narrative.
- **FR-5.2 [V1]** In LLM mode, the system must enforce: a hard character cap per message (to keep WhatsApp messages scannable), a JSON-structured response contract (per-

dimension paragraph + overall paragraph + "lowest dimension" field), and a fallback to Template mode on LLM error, timeout, or schema-violation.

- **FR-5.3 [V1]** The LLM prompt template is itself stored in the config store and editable by super admin. Prompt changes are versioned.
- **FR-5.4 [V1]** Every LLM call is logged with: tenant_id, user_id, session_id, prompt version, input tokens, output tokens, latency, and outcome (success / fallback). PII in prompts is masked in logs.

5.6 Results delivery

- **FR-6.1 [V1]** Results are delivered as a sequence of WhatsApp messages:
  1. One message per dimension (score, max, band, interpretation), in the order defined by the instrument.
  2. One overall summary message (all three dimension scores as a table, overall score, overall band, overall narrative).
  3. One image — the Full-Spectrum Leadership circle — per §5.7.
- **FR-6.2 [V1]** Raw per-question scores are **never** shown to the user. Only section totals, overall total, bands, and interpretation are displayed.
- **FR-6.3 [V1]** All result message templates are tenant-configurable.

5.7 Results visual

- **FR-7.1 [V1]** Generate a personalised circular diagram showing the three dimensions, each segment shaded by band (band → colour mapping configurable per tenant), with the score rendered inside each segment. Output as PNG, attached to the WhatsApp message.
- **FR-7.2 [V1]** If dynamic generation fails or is disabled for a tenant, fall back to a tenant-configured static image plus a text table of scores. The fallback is a feature flag, not a hardcoded path.

5.8 Coaching handoff

- **FR-8.1 [V1]** After results, send the debrief CTA: two tenant-configurable messages offering a coaching conversation with the tenant's assigned coach. Accept `YES` or `NO`.
- **FR-8.2 [V1]** On `YES`, send the coach's booking URL (from tenant config). On `NO`, acknowledge and provide the same URL for future reference.
- **FR-8.3 [V1]** Follow the debrief CTA with a **coaching-interest** question (`YES / NO`). On `YES`, mark the user as a warm lead and trigger a coach notification (§5.9).
- **FR-8.4 [V1]** Close with a LinkedIn / closing message (template, per tenant). After this, the bot sends no further unprompted messages for this session.
- **FR-8.5 [V2]** Real-time AI coaching conversation flow — deferred.

5.9 Notifications

- **FR-9.1 [V1]** When a user flags coaching interest, send the tenant's coach a notification containing: user name, organisation, three dimension scores + bands, overall score + band, lowest dimension, assessment completion timestamp, and a deep link to that user's record in the admin dashboard. Delivery channel (email or WhatsApp) is configurable per tenant.
- **FR-9.2 [V1]** Notifications must be delivered within **5 minutes** of the trigger under normal operating conditions. Failed deliveries retry with exponential backoff up to 3 attempts, then alert the super admin.
- **FR-9.3 [V1]** Coaches can opt in to receive a daily or weekly digest in addition to real-time notifications.

5.10 Voice support

Voice is opt-in per user and can be disabled per tenant via a feature flag.

- **FR-10.1 [V1]** The bot accepts WhatsApp voice notes as input anywhere a text reply is expected. Voice is transcribed via the configured speech-to-text provider, then processed through the same text pipeline.
- **FR-10.2 [V1]** For A/B/C/D answers from voice, a normalisation step extracts the option letter (e.g. "I'd say B", "option B", "bee" → `B`). If extraction is ambiguous, reply with the invalid-answer template and quote the transcript back so the user can confirm.
- **FR-10.3 [V1]** Name and organisation captured via voice are transcribed and echoed back for confirmation (`"I heard 'Priya Sharma' — is that right? YES or NO"`). Any free-text voice input is always confirmed before storage.
- **FR-10.4 [V1]** The user can toggle **voice output** at any time by typing `VOICE ON` / `VOICE OFF` (commands configurable). When voice output is on, long-form messages (welcome, section intros, interpretations, closing) are also sent as TTS voice notes in addition to text. Short messages (single question + options, progress indicators, invalid-answer prompts) remain text-only to avoid noise.
- **FR-10.5 [V1]** STT and TTS providers are configured centrally, pluggable behind an interface, and the selection is logged per call for cost attribution.
- **FR-10.6 [V1]** Transcripts of voice input are stored alongside the text answer for auditability. Audio files are retained per the tenant's data-retention policy (default 30 days, configurable).

5.11 Memory and session state

- **FR-11.1 [V1]** A **session** spans one diagnostic attempt, from first trigger to closing message (or to timeout). Sessions are tagged with `tenant_id`, `user_id`, `instrument_version_id`, `started_at`, `completed_at`, `status`.
- **FR-11.2 [V1]** A **user** is unique per `(tenant_id, whatsapp_phone_hash)`. A returning user (same phone, same tenant) sees a re-entry flow: acknowledgement, then choice to view last results or take the diagnostic again. Re-entry copy is template-driven.
- **FR-11.3 [V1]** Historical results are retrievable — on typing `RESULTS`, the bot re-sends the most recent completed result set.
- **FR-11.4 [V2+]** Longitudinal memory (themes, coaching history, recurring topics) for the AI coaching flow — deferred.

5.12 Escalation and safety

- **FR-12.1 [V1]** The bot is not a counselling or crisis service. If a user's free-text input (during name/org capture or any open prompt) contains language suggesting self-harm, serious distress, or HR/legal issues, the bot must **stop the assessment flow** and send a configurable safe-handoff message directing the user to appropriate resources and the tenant's coach contact.
- **FR-12.2 [V1]** A minimal keyword + LLM-classifier hybrid is acceptable for V1. The classifier threshold and trigger phrase list are configurable.
- **FR-12.3 [V1]** Any escalation event is logged and notifies the super admin and the tenant's coach. User's session is suspended, not deleted.
- **FR-12.4 [V2]** Full contextual safety-model for the coaching flow — deferred.

5.13 Admin dashboard

- **FR-13.1 [V1]** The admin dashboard is a web app with role-based access: `super_admin`, `tenant_admin`, `coach`.
- **FR-13.2 [V1] Content editor** (super admin): CRUD on dimensions, instruments, questions, options, scores, bands, interpretation templates, message templates, LLM prompt templates. All edits are versioned; publishing a new version requires a "publish" action.
- **FR-13.3 [V1] Tenant configurator** (super admin): CRUD on tenants, per-tenant feature flags, coach assignment, branding, WhatsApp routing, instrument selection.
- **FR-13.4 [V1] Interpretation mode toggle** (super admin, per tenant): switch between Template mode and LLM mode. Both modes must be fully configured before LLM mode can be selected.

- **FR-13.5 [V1] User management** (tenant admin, scoped to own tenant): list users, view individual result history, export results (CSV), soft-delete a user (GDPR right-to-erasure).
- **FR-13.6 [V1] Analytics** (§5.14).
- **FR-13.7 [V1] Audit log** (super admin): all content and config edits are logged with actor, timestamp, before/after diff.
- **FR-13.8 [V2]** Knowledge-base upload (company values, strategy, frameworks) — schema reserved in V1, UI deferred.

5.14 Analytics

- **FR-14.1 [V1]** Track, per tenant: sessions started, sessions completed, completion rate, drop-off by question, median time-to-complete, voice-mode adoption rate, coaching-CTA `YES` rate, coaching-interest `YES` rate, booking-link click-through (if trackable via redirect).
- **FR-14.2 [V1]** Track, per tenant: distribution of scores per dimension, distribution of bands, lowest-dimension frequency, overall readiness-level distribution.
- **FR-14.3 [V1]** All analytics honour tenant isolation. Super admin sees cross-tenant rollups.
- **FR-14.4 [V1]** Raw event stream is retained for 18 months (configurable). Aggregates are retained indefinitely.

---

6. Configuration and content model ("nothing hardcoded")

This section defines *what* is in the database and editable, versus *what* is in code. The rule: **anything a non-engineer might reasonably want to change is in the database.**

6.1 In code (not configurable through admin)

- Bot state machine (welcome → name → org → sections → results → CTA → close).
- Scoring arithmetic (sum of option scores; band lookup).
- WhatsApp / STT / TTS / LLM provider integrations.
- Security, auth, tenancy enforcement.

6.2 In the database (editable through admin, versioned)

Entity	What it holds	Who can edit
`dimension`	Name, description, display order, internal tags	Super admin
`instrument`	A named, versioned diagnostic: collection of sections + dimension ordering + bands + defaults	Super admin
`instrument_version`	Immutable snapshot of an instrument at publish-time	System (on publish)
`question`	Stem text, section (dimension), display order, internal tag	Super admin
`option`	Option label (A–D), option text, score value, internal tag, display order	Super admin
`dimension_band`	`(dimension_id, instrument_version_id, min_score, max_score, band_label, interpretation_template)`	Super admin
`overall_band`	`(instrument_version_id, min_score, max_score, band_label, interpretation_template)`	Super admin
`message_template`	Keyed templates (`welcome_1`, `welcome_2`, `section_intro_cognitive`,	Super admin, overridable per tenant

Entity	What it holds	Who can edit
	`invalid_answer`, `abandonment_reminder`, `debrief_cta_1`, etc.) with variable slots	
`llm_prompt_template`	Prompt used for LLM interpretation mode, with variable slots and response schema	Super admin
`tenant`	All tenant config from FR-1.2	Super admin
`coach`	Name, booking URL, notification channel + address, tenants assigned	Super admin
`feature_flag`	Per-tenant toggles (voice, llm_interpretation, email_capture, dynamic_image_gen)	Super admin
`band_colour_map`	Per-tenant or per-instrument band → hex colour mapping for the results circle	Super admin

6.3 Template variable conventions

All user-facing text is a template. Variables use `{{snake_case}}`. The template engine must error on missing variables at render time (fail loud, not silent). Standard variables available in every template:

- `{{name}}` — user's first name
- `{{organisation}}` — user's organisation
- `{{coach_name}}`, `{{coach_booking_url}}`, `{{coach_linkedin_url}}`
- `{{tenant_name}}`
- `{{dimension_names_list}}` — comma-joined list, e.g. "Cognitive Clarity, Relational Influence, Inner Mastery"
- `{{duration_estimate}}` — from instrument metadata
- `{{section_number}}`, `{{section_count}}`, `{{question_number}}`, `{{question_count}}`

Result-only templates also get: `{{score}}`, `{{max_score}}`, `{{band_label}}`, `{{dimension_name}}`, `{{overall_score}}`, `{{overall_max_score}}`, `{{overall_band_label}}`, `{{lowest_dimension_name}}`.

6.4 LLM interpretation mode — prompt contract

The stored prompt template receives the user's full score context plus the tenant's company-context document (V2+; empty string in V1). The model must return JSON conforming to:

```json
{
  "per_dimension": [
    {"dimension": "Cognitive Clarity", "narrative": "..."},
    {"dimension": "Relational Influence", "narrative": "..."},
    {"dimension": "Inner Mastery", "narrative": "..."}
  ],
  "overall_narrative": "...",
  "lowest_dimension": "Inner Mastery"
}
```

Any deviation (missing field, wrong dimension name, narrative exceeding length cap) triggers fallback to Template mode. Length caps are configurable (default: 350 characters per dimension narrative, 500 for overall).

---

7. V1 diagnostic content (Innergy — seeded on day one)

This is the complete V1 diagnostic as it will be seeded into the `instrument` / `question` / `option` / `dimension_band` / `overall_band` tables for the Innergy tenant. Source: `FLS_Diagnostic_for_CXOs`. All scores below are authoritative for V1.

Instrument: `innergy_fls_v1` **Dimensions:** Cognitive Clarity, Relational Influence, Inner Mastery **Total questions:** 25 (8 + 9 + 8) **Maximum total score:** 125 (40 + 45 + 40)

> ⚠ **Scoring reconciliation note** — The source diagnostic document has two internal inconsistencies that are resolved here and must be confirmed by Rashmi before publish:
>
> 1. *Inner Mastery max score*: the question set supports 40 (8 Qs × max 5), but the bands listed in the source are on a 25-point scale. V1 uses **max 40** with proportionally scaled bands (shown below).
> 2. *Inner Mastery Q8*: the source contains two variants. V1 uses the "Leadership Presence" variant (marked *Q8 OPTION 2* in source). The "Grounded Leadership" variant is dropped.
> 3. *Overall score bands*: source overall bands appear to assume a max of 75; V1 uses bands proportionally scaled to the 125-point maximum.

7.1 Section A — Cognitive Clarity (8 questions, max 40)

Section intro (stored as `section_intro_cognitive` template):

> SECTION 1 OF 3: COGNITIVE CLARITY. This section is about how you think and decide — especially when things are fast-moving, ambiguous, or genuinely new. For each question, reply with A, B, C, or D.

#	Stem	Options (A → D)	Scores
A1	When I have to make an important decision but don't have all the information I need, I typically:	A) Wait to see if more information emerges to get clarity before committing · B) Actively find more data to see the full picture · C) Involve team, ask others or escalate upwards · D) Push forward with what I know and adjust as I go	2 · 4 · 2 · 5
A2	When something urgent lands on my desk — a complaint, a request, a crisis — my first instinct is to:	A) Take control, deal with it immediately so it's off my plate · B) Figure out how & why this situation occurred in the first place · C) Delegate it and check back later · D) Pause & see whether it really needs attention and then decide what to do	2 · 1 · 3 · 5
A3	When a decision I made stops producing the results I expected, what below is most representative of you?	A) I find it hard to announce to team that the approach isn't working but eventually I do · B) It is difficult to figure out what the new plan should be · C) It is hard to move fast enough to course-correct before too much damage is done · D) I continue with the plan, sometimes success is just around the corner and does not need a change of direction	4 · 2 · 3 · 1
A4	When someone presents a point of view which challenges my beliefs, I most often:	A) Ask lot of questions to understand the veracity of data · B) Update my view if the evidence is strong enough · C) Feel defensive initially but come around eventually · D) Depends entirely on who is challenging it	4 · 5 · 2 · 1
A5	If I look at how I spend my thinking time in the last month, most of it went toward:	A) Problems and fires that needed immediate resolution · B) Delivering on commitments I already made and need delivery within 1 year · C) Planning and preparing for the next 1–2 years · D) Thinking about where the industry and business are going in 3–5 years	1 · 2 · 3 · 5
A6	When I read about a trend, a technology, or a shift happening in my or another industry, my usual reaction is:	A) I wonder if this could affect us and how · B) I file it away mentally to act on later · C) I actively explore whether there is something here for us · D) I share it with my team to get their perspective	3 · 2 · 5 · 4
A7	When something genuinely unexpected happens — a market shift, a competitive move, an internal crisis — my first internal response is usually:	A) Concern about what this means for my existing plans · B) Curiosity about what this means and what it opens up for me · C) A need to understand it fully before reacting · D) An instinct to protect & double down on what is working right now	2 · 5 · 3 · 2
A8	How quickly do you recognize emerging technological and market shifts before they impact your company or functional strategy?	A) We get to know once it's already impacted us · B) We know and can spot 6–12 months in advance · C) We know and can spot 2–3 years in advance · D) We are looking at weak signals and working on a long-term strategy	1 · 2 · 3 · 4

> Note: A8 has max 4, not 5. Section A max = 5+5+4+5+5+5+5+4 = **38**. This contradicts the source's stated 40. Two options to resolve, for Rashmi to pick: (a) accept 38 as the true max and use the bands below scaled to 38; (b) rewrite A3/A8 so each has a max-5 option. V1 seed below assumes option (a) pending confirmation.

Section A bands (scaled to true max of 38 pending confirmation; original 40-scale bands in brackets):

Band	Range (of 38)	[Original of 40]	Interpretation seed
Strong	31–38	[33–40]	Genuine cognitive agility. Filters well, adapts fast, thinks ahead, stays curious under pressure.
Developing	23–30	[24–32]	Real capability present but inconsistent. Thinks clearly in familiar territory, struggles when situations are genuinely new.
At Risk	14–22	[15–23]	Predominantly reactive and backward-looking. Will be increasingly exposed as the pace of change accelerates.
Critical Gap	8–13	[8–14]	The cognitive infrastructure for leading in complexity is not yet built. Significant and urgent intervention needed.

7.2 Section B — Relational Influence (9 questions, max 45)

Section intro (stored as `section_intro_relational` template):

> Section 1 complete. ✓ SECTION 2 OF 3: RELATIONAL INFLUENCE. This section is about your team — how you lead together, make decisions, give feedback, and respond to change. Answer based on how your top team operates as a whole.

#	Stem	Options (A → D)	Scores	Tag
B1	When my team makes a decision…	A) People comply but may not always fully commit · B) Some decisions get team alignment, but not others — depends on how controversial it is · C) We can consistently converge on a single priority quickly to execute · D) People execute because they genuinely believe in it and not just because it was decided at the top	1 · 2 · 4 · 5	Decision Alignment
B2	When my team is pulled in different directions or facing a complex situation where answer is not clear:	A) We struggle — complexity often leads to confusion or paralysis at the team level · B) We eventually align but it takes longer than it should and creates friction · C) We align reasonably well — someone usually steps up to create clarity · D) We find a way to simplify and align — even when we don't have all the answers	1 · 2 · 4 · 5	Decision Alignment
B3	In my team, how do we treat bad news?	A) Bad news is raised late. By the time we hear it, it is already an established problem · B) It depends on the person · C) Bad news travels up fairly quickly but there are still blind spots we discover late · D) Problems surface early consistently — people feel safe raising things before they are certain	1 · 2 · 4 · 5	Safety
B4	In our team, how do we give feedback?	A) Feedback is process based, people avoid difficult conversations to keep the peace · B) Feedback happens but it is often softened to the point where the message gets lost · C) Feedback is generally direct but occasionally becomes personal or creates defensiveness · D) Feedback is frequent, clear, and the team receives it well without taking it personally	1 · 2 · 3 · 5	Accountability
B5	How does my team operate while working with people with varying levels of experience, personalities and diverse functions and cultures?	A) We experience frequent friction and conflict when working across different teams or cultures and it slows us down · B) We work well within our own team but friction and misalignment tend to surface when we engage across other teams or cultures · C) We generally work well across differences but occasional friction or conflict still surfaces · D) We navigate differences in generation, function, and culture with low friction	1 · 2 · 4 · 5	Cross-functional Influence
B6	When a strategic decision is made, how quickly does our team's operations and execution shift to reflect it?	A) Slowly — it takes significant time before the decision is visible in how we actually work · B) Moderately — we shift eventually but there is always a gap between decision and action · C) Quickly in most cases — the team generally moves fast when direction is clear · D) Very quickly — once a decision is made the team realigns operations almost immediately	1 · 2 · 4 · 5	Responsiveness
B7	When a new priority emerges, how is that treated by our team?	A) We rarely reallocate — existing commitments and budgets tend to stay intact even when priorities shift · B) We reallocate partially — some resources	1 · 2 · 4 · 5	Resource Calibration

#	Stem	Options (A → D)	Scores	Tag
		shift but there is significant resistance to letting go · C) We reallocate meaningfully in most cases · D) We reallocate visibly & decisively — when a new priority is clear the team moves budget, talent and focus		
B8	When a project or initiative is no longer aligned with where we are going, what do we do?	A) We rarely shut things down, projects tend to continue long after they should have stopped · B) We sometimes shut things down but it is slow and politically difficult · C) We generally shut things down when it is amply clear they are no longer working · D) We shut things down quickly and without drama when they no longer serve where we are going	1 · 2 · 4 · 5	Resource Calibration
B9	How frequently do we actively prioritize spotting external "weak signals" to shape our team's strategic discussions?	A) We seek more information to check whether it is the correct read of the market · B) We test quickly & learn · C) We put the finding through our risk management process and make a business case · D) We trust the team who came up with signal and get into action	1 · 2 · 4 · 5	Signal Sensing

Section B max: 9 × 5 = **45**.

Section B bands (as per source):

Band	Range	Interpretation seed
Strong	33–45	This team operates with genuine relational maturity. Decisions land, feedback flows, execution is fast, and the team can let go when needed.
Developing	24–32	Real relational capability exists but is inconsistent. Performs well in familiar conditions, shows friction under complexity or diversity.
At Risk	15–23	The team's relational dynamics are limiting execution and adaptability. Compliance is mistaken for alignment.
Critical Gap	8–14	The team is operating on relational patterns built for a slower, more stable world. Gaps are cultural and require structural intervention.

7.3 Section C — Inner Mastery (8 questions, max 40)

Section intro (stored as `section_intro_inner` template):

> Section 2 complete. ✓ SECTION 3 OF 3: INNER MASTERY. This section is personal. It's about your inner experience — how you feel under pressure, how you recover, and the quality of presence you bring to the people around you. Answer honestly. These questions are often the most revealing.

#	Stem	Options (A → D)	Scores	Tag
C1	When I am in a high-stakes situation — a crisis, a difficult conversation, a major decision under time pressure — my thinking & judgement:	A) Becomes scattered · B) Focus narrows, I could miss wider context · C) Stays clear — pressure doesn't significantly affect my thinking · D) Sharpens — high-stakes situations bring out my clearest thinking	1 · 2 · 4 · 5	Nervous system regulation under pressure
C2	Under a long period of stressful work, my decision-making:	A) I get into fire-fighting mode, more reactive, delay complex decisions · B) I become more cautious & risk averse · C) I stay consistent — I maintain decision quality · D) I remain reliable — I have the tools to maintain clarity	1 · 2 · 4 · 5	Decision Making Under Pressure
C3	When people around me are anxious, stressed, or in conflict, I notice my own inner state:	A) I absorb the emotional atmosphere and it affects my thoughts and productivity · B) I feel the pull but try to maintain my workflow — it takes effort · C) I can feel the tension without it significantly altering my productivity · D) It does not affect me; sometimes I have a positive effect by reducing tension	1 · 2 · 4 · 5	Emotional Stability
C4	After a hard period — a setback, sustained pressure, major conflict — I recover by:	A) Not able to afford recovery time; I snap back and carry it into subsequent situations · B) Recover eventually but it takes more time and deliberate effort · C) I recover with some effort and also carve out time to introspect · D) I recover relatively quickly because I have built practices that help me process and reset	1 · 3 · 3 · 5	Recovery and Reflection
C5	When I step away from work — evenings, weekends, holidays — I:	A) Don't switch off / find it difficult to switch off · B) Can stop working but don't fully feel rested · C) Generally able to disconnect and rest reasonably well · D) Meaningfully disconnect and return more energised than before	1 · 2 · 4 · 5	Rest and Restoration
C6	When I reflect on *why* I lead — the deeper reason beyond the role or results — I find that:	A) I don't often think about this · B) I have a sense of it but it changes depending on circumstances · C) I have a clear enough sense and it helps me stay oriented — especially when things get hard · D) My sense of purpose is what I return to when everything else is uncertain	1 · 2 · 4 · 5	Purpose
C7	When I need to make a judgment call that goes beyond what the data can tell me:	A) I struggle without clear data — I delay, defer, or default to what has worked before · B) I can make the call but don't fully trust it · C) I generally trust my judgment in these moments · D) I am certain in moments which could be technically overwhelming; my clearest thinking often happens in these moments	1 · 2 · 4 · 5	Whole Brain Access
C8	When I think about the effect my presence has on the people around me — my energy, my mood, my level of stress — I would say:	A) I am honestly not very aware of it · B) I am aware when someone points it out or when I notice a visible reaction — but I don't track it proactively · C) I am generally aware and make a conscious effort to manage it, especially in high-stakes situations · D) I have a clear and continuous sense of the energy I am	1 · 2 · 4 · 5	Leadership Presence

#	Stem	Options (A → D)	Scores	Tag
		bringing and actively work to ensure my presence is a resource rather than a drain		

Section C max: 8 × 5 = **40**.

Section C bands (scaled from source's 25-scale bands to the true 40 max):

Band	Range (of 40)	[Original of 25]	Interpretation seed
Strong	32–40	[20–25]	This dimension is a genuine asset in your leadership. The foundation is there. The work now is to extend it, protect it under pressure, and scale it as the organisation grows.
Developing	22–31	[14–19]	Capability exists but is inconsistent. Some situations bring it out, others expose the gap. Targeted investment will create disproportionate return.
At Risk	12–21	[8–13]	A quiet liability — not visible in stable conditions, but under pressure or during change, it shows up as a constraint on performance.
Critical Gap	8–11	[5–7]	Not a development priority — a business risk. Leaders operating here are actively limiting the people around them.

7.4 Overall bands (max 123 pending A-section confirmation; source-scaled)

Assuming Section A max of 38 is accepted: overall max = 38 + 45 + 40 = **123**. If A-section is rewritten to max 40: overall max = **125**.

Overall bands (scaled to 123; adjust linearly on final max):

Band	Range (of 123)	Interpretation seed
AI-Ready	98–123	Your leadership has the foundation to thrive in the AI age. The priority now is staying ahead — building the practices that keep this edge as the pace accelerates.
Partially Ready	72–97	Real strengths with meaningful gaps. In stable conditions this works. As AI raises the bar on speed, judgment, and adaptability, these gaps will become more visible and more costly.
At Risk (Developing Readiness)	48–71	Operating on a model built for a slower world. Gaps across one or more dimensions are significant enough that they are likely already affecting execution, culture, and retention.
Critical (High Risk)	24–47	The leadership infrastructure is not fit for what is coming. Needs a fundamental reset, not another training programme.

8. Data model (reference, not DDL)

Keys are illustrative; actual DDL is the engineer's call.

```
tenant (id, name, whatsapp_number, whatsapp_mode[dedicated|shared],
trigger_payload,
        logo_url, primary_color, secondary_color, linkedin_url, status,
created_at)

coach (id, name, booking_url, notification_channel[email|whatsapp],
       notification_address, created_at)

tenant_coach (tenant_id, coach_id, is_primary)

feature_flag (tenant_id, flag_key, flag_value)

instrument (id, name, description, current_version_id)

instrument_version (id, instrument_id, version_number, published_at,
published_by,
                    metadata_json)  -- metadata includes duration,
dimension_order

tenant_instrument (tenant_id, instrument_version_id, effective_from)

dimension (id, name, description, internal_tag)

section (id, instrument_version_id, dimension_id, display_order,
intro_template_key)

question (id, section_id, stem, display_order, internal_tag)

option (id, question_id, label[A|B|C|D], text, score, display_order,
internal_tag)

dimension_band (id, instrument_version_id, dimension_id, min_score,
max_score,
                band_label, interpretation_template, band_color_hex)

overall_band (id, instrument_version_id, min_score, max_score, band_label,
              interpretation_template)

message_template (id, key, tenant_id[nullable for global default],
                  locale, body, updated_at, updated_by)
-- tenant_id NULL = global default; non-null = tenant override.
-- Lookup precedence: tenant override → global default.

llm_prompt_template (id, key, version, body, response_schema_json,
                     max_narrative_chars, updated_at)

user (id, tenant_id, whatsapp_phone_hash, first_name, organisation,
      voice_mode[off|on], created_at, last_seen_at)

session (id, tenant_id, user_id, instrument_version_id,
status[in_progress|completed|
         abandoned|escalated], started_at, completed_at, abandoned_at)

answer (id, session_id, question_id, option_id, raw_input,
raw_input_type[text|voice],
        voice_transcript_nullable, answered_at)

result (id, session_id, instrument_version_id, cognitive_score,
cognitive_band,
        relational_score, relational_band, inner_score, inner_band,
overall_score,
        overall_band, lowest_dimension_id, interpretation_mode[template|llm],
        interpretation_json, generated_at)

notification (id, tenant_id, coach_id, user_id, session_id,
type[coaching_interest|
              escalation|digest], channel, payload_json, status, attempts,
sent_at)
```

```
audit_log (id, actor_user_id, entity, entity_id,
action[create|update|delete|publish],
           before_json, after_json, at)

event (id, tenant_id, user_id, session_id, event_type, properties_json, at)
-- analytics stream
```

---

9. Non-functional requirements

9.1 Privacy and data protection

- All WhatsApp phone numbers stored as salted hashes; raw number only retained for active session routing.
- Tenant data isolation enforced at the application layer AND via row-level security where the DB supports it.
- User's right to erasure is a soft-delete + scheduled hard-delete within 30 days.
- Audio files retained per tenant's data-retention policy (default 30 days).
- No cross-tenant data access except by super admin, every access logged.

9.2 Security

- Admin dashboard: SSO-capable (SAML/OIDC), MFA required for super admin.
- All secrets (WhatsApp API key, LLM API key, STT/TTS API keys) in a secrets manager, not in env files.
- Rate-limiting on inbound WhatsApp per phone to prevent abuse.
- LLM calls sanitised for prompt-injection (strip user-provided text that resembles system-prompt markers before inclusion).

9.3 Performance

- Bot response to any user message: p95 ⩽ 2s (excluding LLM calls and dynamic image generation).
- LLM interpretation: p95 ⩽ 8s; on exceed, fall back to template mode.
- STT transcription: p95 ⩽ 5s for voice notes ⩽ 60s.
- Image generation: p95 ⩽ 4s; on exceed, fall back to static image.

9.4 Reliability

- 99.5% uptime target for V1.
- Sessions are durable — a backend restart must not lose in-flight sessions.
- WhatsApp webhooks acknowledged within 5s; processing may be async.

9.5 Observability

- Structured logs on every message handled (tenant_id, session_id, user_id hash, event type, latency).
- Metrics: session funnel, voice adoption, LLM cost per tenant, fallback rates, error rates.
- Alerts: notification delivery failures, LLM fallback rate > 10%, WhatsApp webhook lag > 30s.

9.6 Localisation

- V1 ships in English only. Schema supports `locale` on message templates so additional languages can be added without code changes.

---

10. Integration points

Capability	V1 provider (recommended)	Swappable?
WhatsApp messaging	Meta WhatsApp Cloud API (or BSP e.g. Twilio, Gupshup)	Yes, behind `MessagingProvider` interface
LLM (interpretation + future coaching)	Anthropic Claude API	Yes, behind `LLMProvider` interface
Speech-to-text	OpenAI Whisper API or equivalent	Yes
Text-to-speech	ElevenLabs or equivalent	Yes
Image generation (results circle)	Server-side SVG render (e.g. `sharp` + template) → PNG	N/A
Calendar booking	Calendly (URL-based, no integration needed V1)	—
Coach notification	Email (SMTP / SES) or WhatsApp	Configurable per tenant

---

11. Acceptance criteria

A V1 build is considered complete when all of the following pass:

1. **Multi-tenant** — two distinct tenants can be created, each with its own QR code, coach, branding, and instrument selection. A user scanning Tenant A's QR receives Tenant A's experience end-to-end; cross-tenant data leakage is impossible.
2. **Full diagnostic flow** — a user can scan a QR, complete all 25 questions, and receive three dimension messages + overall summary + circle image + debrief CTA + coaching-interest question + closing message.
3. **Scoring correctness** — given a fixed set of answers, the calculated section scores, bands, overall score, and overall band exactly match an independently computed reference (test fixture per §7).
4. **Template mode interpretation** — with LLM mode disabled for a tenant, all result messages render correctly from templates with variable substitution.
5. **LLM mode interpretation** — with LLM mode enabled, results come from the LLM and conform to the JSON schema; forcing an LLM failure triggers fallback to template mode within the latency budget.
6. **Voice input** — a user can answer any question, including name/org capture, via voice note. A/B/C/D extraction works on common voice phrasings.
7. **Voice output** — with `VOICE ON`, long-form messages are delivered as voice notes alongside text.
8. **Admin editing** — a super admin can edit a question, publish a new instrument version, and a new session uses the new version while in-flight sessions continue on the old version.
9. **Notification** — on coaching-interest `YES`, the assigned coach receives a notification within 5 minutes containing all required fields.
10. **Abandonment** — a session left idle for 24h triggers exactly one reminder; a further 48h of silence closes the session.
11. **Re-entry** — a returning user can type `RESULTS` and receive their most recent result set.
12. **Escalation** — input containing configured trigger language halts the session and sends the safe-handoff message + coach/super-admin notification.
13. **Audit** — every content or config edit produces an audit log entry with actor, before, after, and timestamp.
14. **Analytics** — the admin dashboard shows tenant-scoped funnel, score distributions, and CTA conversion rates within 15 minutes of events occurring.

12. Open items for Rashmi / product owner

These are the decisions I could not make from the source docs. Listed in order of how much they block implementation.

1. **Section A scoring max.** Q8 in the source has max 4 (A=1, B=2, C=3, D=4). Do we (a) accept Section A max = 38 and adjust overall max to 123, or (b) add a fifth-point option to Q3 and Q8 to restore max = 40?
2. **Inner Mastery Q8 choice.** Confirm the "Leadership Presence" variant is the intended one (source marked it "Option 2").
3. **Inner Mastery bands.** Confirm the scaled-to-40 bands in §7.3 are acceptable, or supply bands calibrated on the true 40-point scale.
4. **Overall bands.** Confirm the scaled-to-123 bands in §7.4 are acceptable.
5. **LLM mode for Innergy's V1 tenant.** Launch in Template mode or LLM mode? (We recommend Template mode for V1 launch, LLM mode enabled after 2–3 weeks of live data to tune prompts.)
6. **Voice default state.** Should voice be *off by default* (user opts in with `VOICE ON`) or *on by default* (user opts out with `VOICE OFF`)? Recommendation: off by default for V1 — WhatsApp voice notes on every result message can feel excessive.
7. **Escalation trigger list.** We need the seed list of trigger phrases / topics for V1 — this will be a collaboration between product and Rashmi.
8. **Data retention.** Confirm the default audio retention (30 days), session retention (indefinite unless user requests erasure), and analytics event retention (18 months).
9. **Coach notification channel for Innergy.** Email or WhatsApp? What address/number?
10. **Results image — dynamic or static for launch.** Recommendation: ship static fallback on day one, add dynamic generation within the V1 window once core flow is stable.

---

13. Out of scope for V1

- AI coaching conversations (beyond the diagnostic + handoff).
- Knowledge-base upload and RAG grounding (schema is reserved; UI and retrieval deferred to V2).
- Manager / 360 / perception-gap assessments.
- Team aggregate reports.
- In-product payment or coach scheduling (we link out to Calendly).
- Mobile app, web chat, Slack, MS Teams channels.
- Cross-tenant benchmarking.
- Automatic tenant self-signup.

End of PRD v3.0